

Connection Management is HARD

HTTP Workshop 2017, London

It used to be “simple”

Need to make request, ~~open a new connection~~

if an idle connection is open, use that

else, make a new connection (unless there are too many already)



h2: Multiplexing

An h2 connection pretty much always counts as idle


The stream concurrency limit looks like the cap on connections to the same origin

Not really a big change

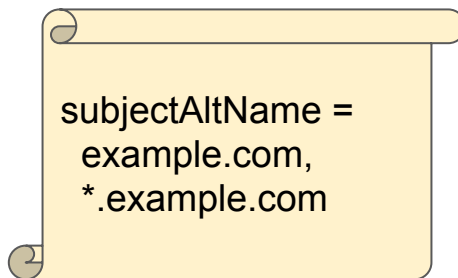


h2: Coalescing

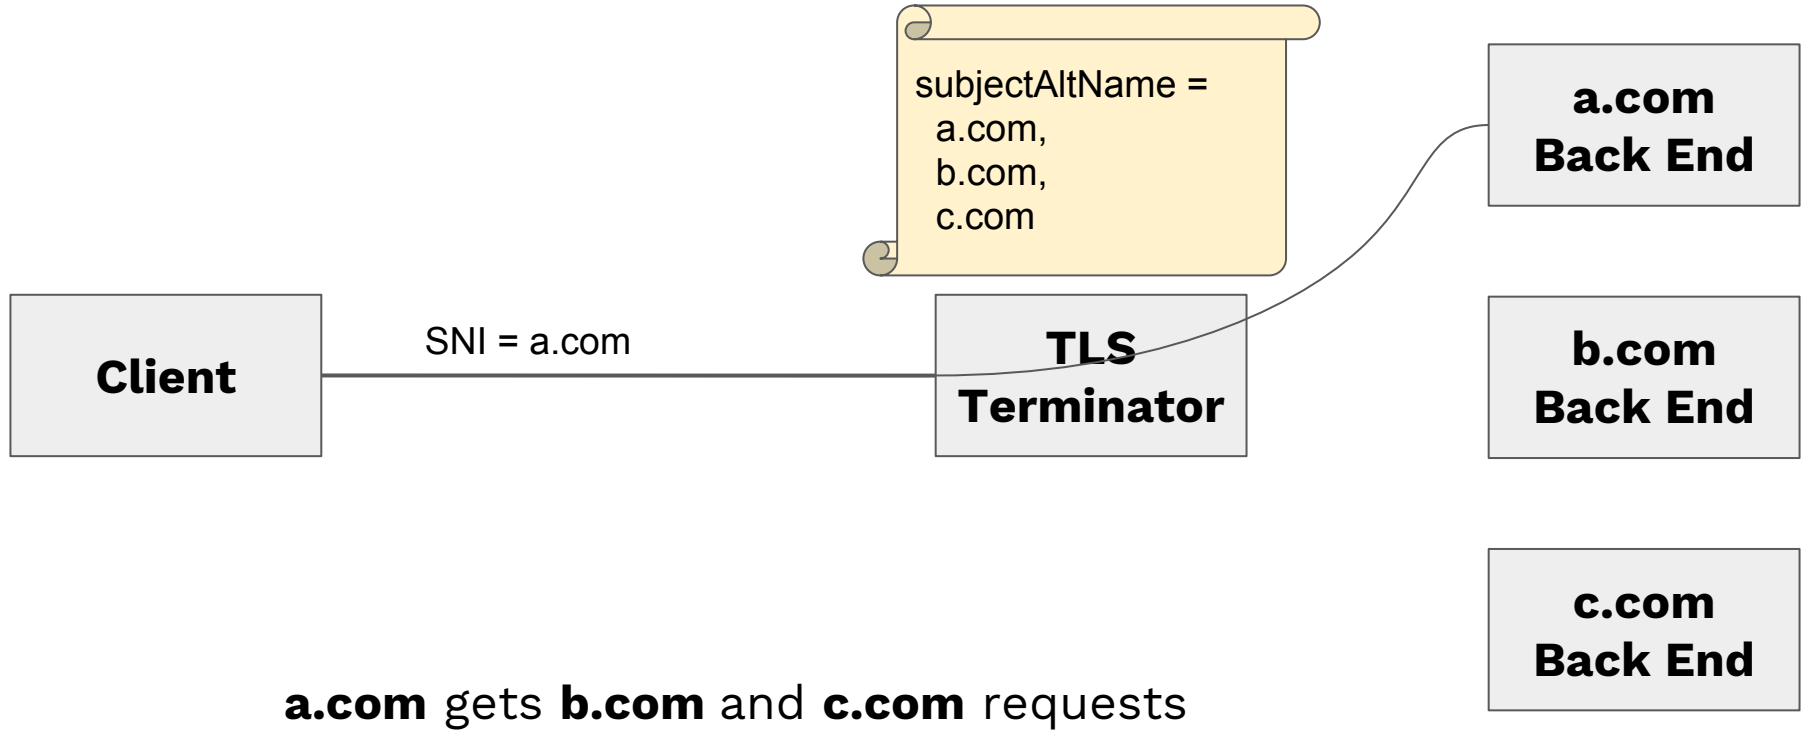
When you need a new connection and the IP address turns out to be the same* as an existing connection, then check the certificate



example.com	192.0.2.75	192.0.2.78	2001:0db8::f2	2001:0db8::f4
www.example.com	192.0.2.75		2001:0db8::f2	2001:0db8::f4




Coalescing -> Operational Challenges



421 (Not Authoritative)

We added 421 to deal with this problem, but it's a little slow



ORIGIN Frame


A turgid history of changes

First it was simple: a list of origins

It grew statement management: additions, removals, flushes

A minimalist “nothing other than the SNI” option was considered

But ended up closer to the initial design



ORIGIN Frame Today

Initially a connection is valid for any host in the certificate (h2 default)

An ORIGIN frame indicates support for the feature

The set of origins is then

- a single origin (https, SNI hostname, server port**)

- plus origins that are included in ORIGIN frames

- less origins that generate responses with a 421 status code

Alt-Svc

Used to prime cross-protocol upgrades (opp-sec, QUIC)

Alt-Svc: h2="x.example.net:443"

= Connect to x.example.net instead

Alt-Svc: h2=":8443"

= Connect to port 8443 instead

Alt-Svc: hq=":443"

= Use QUIC

Alt-Svc

CNAME at the HTTP layer

Doesn't affect authority

Changes ALPN protocol, DNS lookup, and (TCP) port

Make before break

An invisible transition is great because users don't notice

which is **terrible** because it's hard to see when the alternative breaks

Secondary Certificates

h2 allows coalescing, but everything has to fit in one certificate

For example: *.www.yahoo.com, add.my.yahoo.com, *.att.yahoo.com, att.yahoo.com, au.yahoo.com, be.yahoo.com, brb.yahoo.com, br.yahoo.com, ca.my.yahoo.com, ca.rogers.yahoo.com, ca.yahoo.com, ddl.fp.yahoo.com, de.yahoo.com, en-maktoob.yahoo.com, espanol.yahoo.com, es.yahoo.com, fr-be.yahoo.com, fr-ca.rogers.yahoo.com, frontier.yahoo.com, fr.yahoo.com, gr.yahoo.com, hk.yahoo.com, hsr.d.yahoo.com, ideanetsetter.yahoo.com, id.yahoo.com, ie.yahoo.com, in.yahoo.com, it.yahoo.com, maktoob.yahoo.com, malaysia.yahoo.com, my.yahoo.com, nz.yahoo.com, ph.yahoo.com, qc.yahoo.com, ro.yahoo.com, se.yahoo.com, sg.yahoo.com, tw.yahoo.com, uk.yahoo.com, us.yahoo.com, verizon.yahoo.com, vn.yahoo.com, www.yahoo.com, yahoo.com, za.yahoo.com, zed.yahoo.com

Idea: Add a Certificate to a Connection

This would allow a server to have many small certificates

certificates can be added as needed

Great for servers that have multiple names

Especially good for servers with multiple disparate names

This could also fix or help with

post-handshake client authentication

domain fronting for privacy reasons

How it works

TLS exported authenticators copy the authentication messages from TLS into a blob that can be used in a higher layer protocol

On Stream 0:

- An endpoint can send an authenticator with CERTIFICATE

- An endpoint can request a CERTIFICATE using CERTIFICATE_REQUEST

After make or receiving a request:

- An endpoint can reference a certificate with USE_CERTIFICATE

- Request USE_CERTIFICATE by sending CERTIFICATE_NEEDED

Examples

